

	Lustre	NFS	Others	Description	Reference
hero node	cluster		process (distributed)		
client server					
message					
file					
file-set	file-set (planned)				
component object (cob) /	stripe object	component data file		metadata-enhanced storage object	<a href="https://docs.google.com/a/x">https://docs.google.com/a/x</a>
	import, export	session			
	plrpc connection	connection			
		channel			
reintegration stream	last rcvd	reply-cache			
	[ad client data	(session, slot) pair			
	client mid	client-owner		a client identifier stable across client restarts	
	client uuid	client-id			
fid	fid	file handle (roughly)	oid		
inode number	client inode number	fileid			
layout	LOVEA	layout			
	CMD	meta-data striping			
operation					
update					
node-id	mount	mount			
	cluster				
dependency of operations					
dependency of updates	DLM				
	grant				
user					
credentials					
capability	recovery		a lock (as opposed to a latch, NOTE: it would be great to have separate terms for a "lock request"		
	DLM lock				
	lock enqueue				
	lock cancel				
	lock invalidation				
	blocking AST				
	completion AST				
	lock owner				
	busy lock (address)				
				a file attribute recording a name and a parent directory of the file. In the current data-model, back-pointers are maintained as a part of object-index.	
back-pointer	(planned)	---			
epoch	epoch				
replay	replay				
resend	resend				
file version	inode version	change attribute			
	xid	sequence-id			
	MDS				
	MDT	NFS server			
	MGS			management server	
	OS				
	OST	NFS server			
persistent (of an operation or an update)					
stable (of an operation or an update)					
	extent lock				
	---			whole-file lock given out by a meta-data server	
	---	---	escrow locking (data-bases community)	escrow lock on a file attribute	
					<a href="https://docs.google.com/a/horizontalstorage.com/doc?id=0AaglcQbR4ems2GhYV2hqdmdfMTIwZjZlOGY5ZjZl&amp;hl=en">https://docs.google.com/a/horizontalstorage.com/doc?id=0AaglcQbR4ems2GhYV2hqdmdfMTIwZjZlOGY5ZjZl&amp;hl=en</a>
FOL	changelog (not really)	---	journal (local file systems), log (databases community)	File Operation Log: an ordered list of descriptions of executed operations maintained by a MO back-end	
FOL record	---	---	record, entry		
					<a href="https://docs.google.com/a/horizontalstorage.com/doc?id=0AaglcQbR4ems2GhYV2hqdmdfNjE1Y3Q5bXU2GQ&amp;hl=en">https://docs.google.com/a/horizontalstorage.com/doc?id=0AaglcQbR4ems2GhYV2hqdmdfNjE1Y3Q5bXU2GQ&amp;hl=en</a>
FOP	RPC	procedure, operation		File Operation Packet	
ADDB	---	---		Analysis and Diagnostics DataBase	
ADDB record	---	LAYOUT/COMMIT (roughly)			
ATOM	SOM			ATribute On Meta-data server	
basic file attributes		stat, data (reiserfs)		file attributes required to perform a stat(2) call	
NRS	NRS	---		Network Request Scheduler	
orphan	orphan				
durability notification (of an operation or an update)	last committed				
		directory delegation			
		open delegation			
	epoch boundary, cut	---	consistent cut, consistent snapshot (distributed algorithms community)		
meta-data intent	WBC	compound (roughly)	meta-data cache		
data intent	---				
	hash	readdir cookie			
cached IO					
nocache IO					
direct-IO					
fsync	fsync	fsync	fsync (POSIX)	read or write calls	
IO	IO	IO	IO	network data transfer	
transfer	IO, transfer (CLIO)	IO	IO	data transfer to or from persistent storage	
reintegration	reintegration		merging		
					<a href="https://docs.google.com/a/horizontalstorage.com/doc?id=0AaglcQbR4ems2GhYV2hqdmdfNjE1Y3Q5bXU2GQ&amp;hl=en">https://docs.google.com/a/horizontalstorage.com/doc?id=0AaglcQbR4ems2GhYV2hqdmdfNjE1Y3Q5bXU2GQ&amp;hl=en</a>
	eviction	drop expired lease			
undo	undo		undo (data-bases community)		
redo	redo		redo (data-bases community)		
redo-tail	---	---	---		
storage application			plugin		
back-end					
cache pressure			memory pressure, working-set balancing, cache trimming (kernel VMs)		
			VMI-wireout, laundromat, page stealing thread, pageout daemon (kernel VMs)		
cache pressure handler	state machine (CLIO)	---	state machine		
snapshot			snapshot		
clone			clone (OSD-2)		
networking transport					
COW	---	---	COW	Copy-on-write	
			in-place update (data-bases community, local file systems), home update		
fid home server	fid home server		(local) transaction		
			transaction group		
			transaction open		
			transaction close		
			transaction commit		
commit call-back	commit call-back	---	commit call-back		
	block number		block number, disk address		
data	data	data	meta-data		
meta-data	meta-data	meta-data	meta-data		
			data or meta-data		
operation error				semantic operation error: rename of a non-existent file, wrong permissions, memory allocation failure, etc., Guarantees: file system consistency is maintained, File system state is consistent with that reported to applications, node, storage or network failure leading to recovery, Guarantees: consistency of visible file system state is maintained, File system state is eventually consistent with that reported to applications.	
failure				node, storage or network failure breaking system model assumptions, for example, more failures than a given RAID mode can recover from, or a loss of all redundant copies of meta-data, Guarantees: consistency of visible file system state is maintained, Some state reported to applications can be no longer accessible, Clients holding no longer relevant state are (optionally?) evicted.	
catastrophe				violation of internal system invariants, or unrecoverable failure during recovery from some other failure, Guarantees: none,	
system failure					
parity group		parity group (RAID literature)		a collection of data blocks (data units) sharing parity	
SNS	SNS (planned)	RAID5, network RAID		server network striping	
CSIM				Cluster storage system simulator, uses ADDB	
MO				ClusterStor Colibri Storage Software, aka Colibri, M0	
DS				Data Server node (can be proxy, cache, or final server)	
FDML				File Data Management Language	
C4				ClusterStor Colibri Control Center, backend management server plus BUI, based on Trinity	
extent				continuous subspace of linear namespace	
device				Terms with the word "device" will refer to notions associated with a single physical storage device.	
				a disk, flash card, raid array, memory or similar device note that some physical storage devices offer persistence, others do not, C4 is responsible for uniquely labeling a physical storage device	
physical storage device				a physical storage device attached to a M0 node	
attached storage device				metadata store which is part of a storage device. A metadata store provides a variety of indexes	
metadata container				linear address space which is part of a storage device. The data stored in a data container can be accessed through the linear offset or through objects described in a metadata container,	
data container					<a href="https://docs.google.com/a/x/yratex.com/Edit?docid=QATg1HfU2caZ2QhNWXn4oXpM4aXNwTJkZG9">https://docs.google.com/a/x/yratex.com/Edit?docid=QATg1HfU2caZ2QhNWXn4oXpM4aXNwTJkZG9</a>
container				a data container or a meta-data container	
				A Container Object stores other objects inside. A storage device can hold	<a href="https://docs.google.com/a/x/yratex.com/Edit?docid=QATg1HfU2caZ2QhNWXn4oXpM4aXNwTJkZG9">https://docs.google.com/a/x/yratex.com/Edit?docid=QATg1HfU2caZ2QhNWXn4oXpM4aXNwTJkZG9</a>